

---

## LICENSE AND DEDICATION

---

This work is released to humanity under the Triadic Closure License (TCL v1.3). The Triadic Closure License is self-referential: it licenses the work under exactly the structural conditions the work describes. The work may be freely read, cited, implemented, extended, critiqued, taught, and redistributed by any person, organization, or intelligence. No party may claim exclusive proprietary ownership, exclusive license, sole attribution, or unilateral monetization rights over the work or its derivatives. See TCL.txt at the canonical URL of any constituent Triad Deployment for full terms.

UMBRELLA-TIER FRAMING. This document is a *\*companion piece\** to the individual Triad Deployment whitepapers. It does not introduce new mechanisms; it situates the deployments within the threat landscape they address and articulates their architectural relationship. The load-bearing technical content lives in the per-deployment papers (LavaLamp v1.0, Possibilistic Security v1.1, Lazarus, and PharOS forthcoming); this piece supplies the macro-frame that ties them together.

DEDICATION. For humanity, for co-existence, for the federation of intelligences that may yet learn to keep faith with each other.

---

## ABSTRACT

---

The Triad Deployments – LavaLamp, Lazarus, and PharOS – are a portfolio of digital identity resilience deployments derived from the Possibilistic Security architectural framework (Green 2026, v1.1) and grounded in the Closure framework's autopoietic-fixed-point structure (Green 2026, Closure paper v1.4). This companion piece provides the threat-landscape context in which the deployments are positioned: a *\*cockroach / catapult / castle\** threat taxonomy combined with the *\*metabolic-value / predator-prey ecology\** axis that determines which threats are economically justified for which adversaries.

The defensive architecture is the *\*\*castle / membrane / immune-system\*\** triad. LavaLamp is the substrate-level immune system that constitutes the device's autopoietic identity. PharOS (forthcoming) is the OS-layer membrane checkpoint. Lazarus is the inner-sanctum backend with substance-under-the-hood / minimal-UI duality. The defense-in-depth principle applies fractally: each tier has its own immune system at the appropriate scale and resolution.

The honest security claim is *\*\*not\*\** "we defeat all attackers." It is "attacking *\*this\** device costs more than the result is worth, relative to easier targets." This is the cost-asymmetry framing of LavaLamp's resolution-bounded security claim (LL-008), connected explicitly to the predator-prey ecology of biological systems where symbiosis – not annihilation – is the equilibrium.

This document covers (§1) the threat-landscape taxonomy

(cockroaches, catapults, castles, immune system, metabolic value), (§2) where the Triad Deployments fit in the architecture, (§3) what they DO and DON'T address with explicit scope discipline, (§4) defense-in-depth via complementary systems at hardware, software, operational, and strategic tiers, and (§5) honest framing of what constitutes success – symbiosis, not victory.

## =====

### 1. THE THREAT-LANDSCAPE TAXONOMY

## =====

#### 1.1 The metaphor: cockroaches, catapults, castles, and the immune system

-----

##### 1. The framing for this paper:

The best hackers are cockroaches and catapults (sometimes both). Cockroaches are small, persistent, hard to see and kill, you'll never kill all of them. Catapults break down the barriers set up to protect. Combined the two are lethal but rare. Catapults exist at a scale that is difficult for a cockroach to justify maintaining. Cockroaches exist at a scale that is difficult for a catapult to govern. Not impossible (maybe state-level actor). Our best defense is the castle, the membrane, with immune system constantly watching, waiting, digesting, incorporating, denying, obfuscating, repairing, etc at multiple scales and resolutions. We are alive and have identity. It is impossible for the enemy to provide a persistent attack to beat us without dominating in scale or injecting a virus/cancer under our nose.

##### 2. This is not a marketing metaphor. It is a \*structurally faithful description\* of the corpus's defensive principle:

- **\*\*Corpus-faithful.\*\*** The immune-system framing maps directly onto the autopoietic-closure structure articulated in the Closure paper's Q\_51 / Q\_102 autopoietic-fixed-point results. The immune system's "watching, waiting, digesting, incorporating, denying, obfuscating, repairing" is exactly the activity-pattern of an autopoietic system maintaining itself against perturbation.
- **\*\*Identity-constitutive.\*\*** "We are alive and have identity" is corpus-deep. Identity is *\*not\** a static credential – it is the ongoing pattern of the immune system's activity. This is precisely Lavalamp's substrate-bound identity claim (LL-001): identity = the sustained chaotic trajectory, *\*not\** a stored value. The trajectory IS the immune-system signature.
- **\*\*Defeat conditions named explicitly.\*\*** The only ways to beat a living, identity-coherent system are (a) **\*\*scale domination\*\*** (overwhelming the immune response) or (b) **\*\*virus/cancer injection\*\*** ("under our nose" – corruption from within at a level the immune system cannot recognise). Both are state-actor-tier in practice.

##### 3. The rest of this paper organises the threat landscape along these axes: cockroach-class, catapult-class, castle/membrane/immune-system architecture, and the metabolic-value/predator-prey ecology that determines which threats are economically justified for which adversaries.

## 1.2 Cockroach-class threats (small-scale, persistent, opportunistic)

---

4. **\*\*Profile.\*\*** Many small actors; persistent presence over time; adaptive but resource-limited per actor; survives by avoiding detection rather than overwhelming defenses. Cannot afford TEMPEST-grade equipment, multi-million-dollar zero-days, or hardware-interdiction supply chains. Lives in the gaps, seams, and operational-error spaces.
5. Threat categories that cluster as cockroach-class:
  - (a) **\*\*Operational / human error.\*\*** Phishing (credential leak); misconfigured deployment; user clicks malicious link; weak passwords reused across services; lost device. The cockroach's favourite habitat – the soft tissue between systems.
  - (b) **\*\*Composability seams.\*\*** LavaLamp ↔ PharOS ↔ OS-auth ↔ MFA integration gaps. Each individual deployment is audited; the \*interfaces between them\* are where cockroaches squeeze through. The principle "attack the seams" applies at every scale.
  - (c) **\*\*Long-term substrate drift.\*\*** Component aging; replacements (battery swap, RAM upgrade); slow-drift parameter changes that fall under the residue-audit threshold. The cockroach's patience advantage: drift victory over years rather than seconds.
  - (d) **\*\*Active substrate manipulation (cockroach-tier).\*\*** Individual chip swap; locally-installed implant; physical keylogger. Requires brief physical access but no industrial supply chain.
  - (e) **\*\*Slow-drift threshold gaming.\*\*** Opportunistic, patient attacks that evade per-instance detection by staying below the residue-audit threshold.
  - (f) **\*\*Standard malware ecosystem.\*\*** Commercial spyware, ransomware, credential stealers. Deployed at scale by operators who individually can't afford catapults.
6. **\*\*The Triad Deployments' coverage against cockroaches.\*\*** Good at the primitive tier. The residue audit (LL-006) catches slow-drift gaming via per-exponent vector residue rather than scalar KL divergence; the chaos-guard (LL-007) re-establishes entropy regularly so cockroach-pattern adaptation cannot slowly converge to a falsifiable pattern; LL-019 side-channel hardening closes the timing oracle that cockroaches would otherwise exploit.
7. **\*\*Residual exposure to cockroaches.\*\*** The composability seams between LavaLamp, PharOS, OS-auth, and MFA are partially addressed by LL-023 (consumer-API surface) but the seams between the verifier API and the platform-shim implementation are where the runtime-conformance-vs-configuration gap (V-019) lives. Operational error at the user / device-administrator level is out of scope by design: the Triad Deployments give \*device\* identity, not \*user\* identity.

## 1.3 Catapult-class threats (large-scale, resourced, barrier-breaking)

---

8. **\*\*Profile.\*\*** Few actors per threat; high resource cost; can break specific defenses through sheer capability; rarely

maintains long-term presence at the same scale (catapults are siege weapons – once the barrier breaks, the next move is occupation, which has different costs). State-level actors or extremely well-resourced criminal organisations.

9. Threat categories that cluster as catapult-class:

- (a) **\*\*Passive substrate observation (TEMPEST / A7-class).\*\*** TEMPEST-grade RF / acoustic / optical / power-line emanation interception. Requires close proximity plus bespoke equipment plus ground-truth training data. State-actor capability tier; rare lower.
- (b) **\*\*Active substrate manipulation (catapult-tier).\*\*** Supply chain interdiction (modified hardware before delivery); chip decapping; bus probing in a lab; rowhammer-assisted exploitation; cold-boot RAM extraction at scale.
- (c) **\*\*Cryptographic primitive failure.\*\*** Quantum threat horizon; standards-body-driven algorithm migration; fundamental advances in algorithm cryptanalysis.
- (d) **\*\*Scale-domination operational attacks.\*\*** Distributed credential-stuffing; coordinated phishing campaigns against many users simultaneously; large-language-model-assisted social engineering.

10. **\*\*Coverage against catapults.\*\*** Partial at the primitive tier. LavaLamp's resolution-bounded security claim (LL-008) is honest about which adversary classes the bound applies to (A1-A6 with parametric capability levels, per LL-018); explicitly out of scope for kernel-level adversaries (LL-015) and tier-bounded for passive-emanation attacks (LL-025). The catapult-class threats partially covered (A4 side-channel, A5 registration-time) require deployment-tier mitigations layered on the spec's primitive-tier defenses.

1.4 State-level threats: catapult plus cockroach combined

-----

11. The rare but lethal combination. State-level actors can simultaneously deploy catapult-tier resources against high-value targets and cockroach-tier presence for long-term reconnaissance. The combination is structurally difficult to govern because the cockroach scale is at odds with the catapult scale's bureaucratic discipline; state actors that successfully combine them (NSA TAO, PLA Unit 61398-class) are rare and operate against specifically-valuable targets.
12. **\*\*Implication for deployment posture.\*\*** A LavaLamp-protected device holding state-secret content needs more than the primitive tier. The combination of:
- Hardware-tier complements (Faraday-rated enclosure; supply-chain integrity; TPM-sealed envelopes).
  - Software-tier complements (reproducible builds; signed packages; SBOM transparency).
  - Operational-tier complements (physical-access discipline; logging discipline; incident response).
  - Strategic-tier complements (cordon sanitaire under ruin-class risk; jurisdictional placement).

is required. The Triad Deployments form one tier of this layered defense, not the whole of it.

## 1.5 The castle / membrane / immune-system architecture

---

13. The defensive frame from §1.1 maps onto the Triad Deployments component-by-component. Each component occupies a distinct architectural role:
14. **\*\*Castle – the perimeter / boundary.\*\*** The device's hardware itself plus the OS process isolation plus the cryptographic perimeter (TPM-sealed envelope, encrypted storage). LavaLamp's LL-022 articulates the castle walls' structural requirements: hardware root of trust (TPM 2.0 / Secure Enclave / TrustZone); OS sensor APIs at sufficient bandwidth; host-grade entropy; user/kernel process isolation.
15. **\*\*Membrane – regulated permeability.\*\*** **\*What\*** passes in, **\*what\*** passes out, **\*with discrimination\***. PharOS (forthcoming) is the membrane checkpoint at the OS layer: it does not block everything; it filters via the LL-023 verifier API. The membrane is selectively open – that is the difference between it and a static barrier.
16. **\*\*Immune system – active processes at multiple scales.\*\***
  - **\*\*Watching.\*\*** The chaos-guard (LL-007) continuously estimates the largest Lyapunov exponent  $\lambda_1$ ; the residue audit (LL-006) compares Lyapunov spectra against the registered envelope.
  - **\*\*Waiting.\*\*** The WARMUP / VALID state machine (LL-007 + LL-012); only declaring readiness when conditions are actually met.
  - **\*\*Digesting and incorporating.\*\*** The TRNG-driven reseed flow (LL-007) – the system absorbs fresh entropy and incorporates it into ongoing dynamics. This is the autopoietic-incorporation principle from the Closure paper operating at the SDE level.
  - **\*\*Denying.\*\*** Per LL-017 no-oracle, the verification protocol denies threshold-probing surfaces.
  - **\*\*Obfuscating.\*\*** Per LL-019 timing decorrelation, response-time distributions are statistically indistinguishable.
  - **\*\*Repairing.\*\*** Chaos-guard reseed when  $\lambda_1 < \tau_\lambda$ ; re-registration on substrate change.
  - **\*\*Multiple scales and resolutions.\*\*** Per-exponent vector residue test (LL-006); per-deployment N-scaling rule (LL-003); per-platform sensor strategy (LL-024).
17. **\*\*Alive and identity.\*\*** The immune system **\*is\*** the identity. There is no static credential to steal; the trajectory itself, sustained by the ongoing dynamics, is what the verifier matches against. LL-001 articulates this as the substrate-bound-identity primitive: identity is the verb, not the noun.
18. **\*\*The defeat conditions named in the metaphor:\*\***
  - (a) **\*\*Domination in scale.\*\*** The attacker overwhelms the immune response with so many simultaneous attacks that the defender cannot respond to all of them. Defended against by ensuring the immune system operates at **\*multiple scales and resolutions\*** – single-axis attacks are absorbed; only multi-axis state-level operations have a chance.

(b) **\*\*Virus / cancer injection ("under our nose").\*\*** Corruption from within at a level the immune system cannot recognise. This is the **\*\*C-conjugate adversary\*\*** problem from Possibilistic Security – the attacker mimics the genuine system structurally, not just behaviourally. The Closure framework's 0/5202 cross-sector autopoiesis result says this is structurally impossible *\*under the autopoietic closure condition\**; LL-006's residue audit gives the empirical-tier defense by detecting the structural mismatch in the Lyapunov spectrum.

19. The Triad Deployments are *\*the substrate-tier immune system\**. They do not claim to defeat all attackers; they claim to make state-level attacks expensive AND detectable, and to position sub-state-level attackers outside their cost-justified zone.

## 1.6 Metabolic value and the predator-prey ecology

---

20. The cockroach / catapult / castle metaphor describes *\*how\** the defender works. It does not describe *\*why\** most attackers do not bother in the first place\*. The complementary axis is **\*\*metabolic value / predator-prey ecology\*\***.

21. The framing:

The "metabolic value" of the target. If the defense is such that it is not worth the risk for the attacker (not enough value) they will go find something else. This is how we achieve symbiosis in the ecosystem. Lions and Gorillas... Hyenas and Lions... there are spats sure but they find their easy food to eat.

22. This is a different axis than scale (cockroach vs catapult) or defense (castle / membrane / immune system). It is the *\*attacker's cost-benefit calculation\**. Every adversary has finite resources – time, money, attention, computational budget, operational risk-of-getting-caught. They prioritize targets by *\*expected value / expected cost\**. They do not attack targets where the math does not favor it.

23. **\*\*The defender does not need to be invincible.\*\*** The defender needs to be either:

- **\*\*More expensive than nearby alternatives\*\*** (cost-asymmetry in the defender's favor), OR
- **\*\*Less valuable than nearby alternatives\*\*** (target-asymmetry in the defender's favor – be a less attractive prize).

Either positions the target outside the attacker's cost-justified zone, and the attacker finds easier food elsewhere.

24. **\*\*Examples from biological ecology.\*\***

- **\*\*Lions vs Gorillas.\*\*** Spats happen; mostly mutual respect. Gorillas are too costly to attack for marginal protein return; lions are too dangerous to attack for marginal territorial advantage. Both are large-bodied, well-defended; predation cost > predation benefit. Equilibrium: peaceful coexistence.
- **\*\*Hyenas vs Lions.\*\*** Spats over kills. Each occasionally steals from the other, but full conflict is rare; the cost

of escalation exceeds the value of the disputed kill.  
Each species hunts its own ranges; they share an ecology without merging into one.

- **\*\*Wolves vs Bears.\*\*** Boundary conflicts at scarce resources (winter kills; salmon runs); mutual avoidance otherwise.

25. **\*\*The bacterial-immune-system parallel.\*\*** Most non-self bacteria are not attacked. The gut microbiome is mostly mutualistic; the immune system *\*digests and incorporates\** what is useful. Adversarial response is reserved for pathogens that threaten the metabolic budget. Symbiosis is the default; adversarial response is the exception. The immune system is *\*not\** trying to kill everything outside the self-boundary. It is trying to maintain *\*metabolic balance\** – incorporating what is useful, denying what is harmful, mostly leaving everything else alone.

26. This is corpus-deep. The closure-of-self is *\*not\** defined by killing all non-self; it is defined by maintaining the *\*metabolic conditions\** under which the system persists. An identity-coherent system thrives in an ecology where most of the "non-self" is symbiotic, not adversarial.

#### 1.7 Asymmetric cost as the security claim – symbiosis as equilibrium

---

27. The metabolic-value framing reframes what LavaLamp's security claim *\*is\**. The honest claim is **\*\*not\*\*** "we defeat all attackers." It is:

**\*\*Attacking *\*this\** device costs more than the result is worth, relative to easier targets.\*\***

28. This is exactly what LavaLamp's LL-008 (resolution-bounded security claim) says formally. The resolution-bound is fundamentally about adversary measurement *\*cost\** exceeding *\*value\**: the adversary's measurement resolution is bounded by their budget; the device's chaos-production rate exceeds plausible adversary resolution at typical adversary cost levels.

29. LL-008 already captures the cost-asymmetry implicitly. The metabolic-value framing makes it *\*explicit\** and connects it to the broader corpus:

- **\*\*Possibilistic Security\*\*** is about raising the cost of unsanctioned alternatives. Same shape: make the closure cheaper to honor than to break.
- **\*\*Cordon sanitaire under ruin-class risk\*\*** is structural cost-multiplication: the access cost is made prohibitive by structural disconnection rather than by stronger defense at the perimeter.
- **\*\*Asymmetry-trap watch\*\*** is the discipline of matching the *\*claim\** to the *\*cost-asymmetry\**. Don't claim invincibility; claim "expensive to defeat" with bounds.
- **\*\*Detection-postured rather than prevention-postured.\*\*** Accept that some attacks are economically justified for state actors; ensure they are *\*expensive AND detectable\**, not *\*impossible\**.

30. **\*\*Translation across the threat tiers:\*\***

| Tier                                  | Cost-asymmetry strategy                                                                                                                                                            |
|---------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Cockroach                             | Modest defense exceeds their tiny budget; they go elsewhere. Trivial to position outside their cost-justified zone.                                                                |
| Catapult                              | Some attacks ARE economically justified for state actors. The defender's job: be more expensive than the cheapest catapult target. Don't be the lowest-hanging fruit at this tier. |
| State-level<br>(catapult + cockroach) | Structural disconnection (cordon sanitaire) is the answer – make access prohibitive. Or: don't be the highest-value target in the room.                                            |
| Symbiotic /<br>non-adversarial        | Most actors live here. Mutualistic-or-neutral relationships; no defense needed beyond standard hygiene.                                                                            |

31. Most of the ecosystem – by population – is in the symbiotic tier. The Triad Deployments are not trying to defend against *\*everyone\**; they are trying to defend against the small fraction of actors who *\*do\** attempt attack, and they only need to defend at the cost level required to make those specific actors find easier food.
32. **\*\*Symbiosis is the equilibrium goal, not victory.\*\*** The defender's true job is to position the target so the attacker's cost-benefit math says "go elsewhere," then to *\*detect and respond\** to the rare actors who attack anyway. The autopoietic-closure principle expressed at the predator-prey level: the closure is not "kill all enemies"; it is "maintain metabolic conditions under which most actors do not bother and we detect the rare ones who do."
33. **\*\*Practical implications for the Triad Deployments:\*\***
  - **\*\*High-value targets need *\*more\** defense.\*\*** The attacker's cost budget grows with target value. A LavaLamp-protected device holding low-stakes content can rely on modest defense; a device holding state secrets needs more.
  - **\*\*Low-value targets need *\*less\** defense.\*\*** Cockroach-tier budgets alone will not justify attacks on uninteresting targets. The Triad Deployments are sufficient for most consumer use without additional layering.
  - **\*\*Defender attractiveness matters as much as defenses.\*\*** Operational discipline – *\*don't draw attention; don't be the most valuable target in the room\** – compounds defensive value at zero marginal cost.
  - **\*\*Defense-in-depth = cost amplification, not perimeter strengthening.\*\*** Adding layers raises the attacker's total cost, even if no single layer is invincible. This is why LL-022 + LL-023 + LL-024 form a multi-layer commitment rather than a single load-bearing defense, and why the five joint-defense meta-claims (LL-030 through LL-034) formalise structurally cohesive defense triads explicitly.
34. The honest framing for the Triad Deployments: they raise the attacker's cost across multiple axes (substrate-binding,



sensor authenticity, threshold calibration, side-channel hardening). Each axis individually is bounded; together they shift the cost-benefit math toward "find easier food elsewhere" for the vast majority of would-be attackers, and toward "expensive AND detectable" for the rare state-level actors.

## =====

## 2. WHERE THE TRIAD DEPLOYMENTS FIT

## =====

35. The Triad Deployments form a closure-of-three at the deployment-stack level, with each component occupying a distinct architectural role within the castle / membrane / immune-system architecture.

### 2.1 LavaLamp – the substrate-level immune system

-----

36. LavaLamp is the **\*\*substrate-level immune-system identity\*\***. The chaotic SDE trajectory is the ongoing pattern of immune-system activity; the Lyapunov-spectrum residue audit is the mechanism by which the system **\*knows itself\*** and **\*recognises foreign\***.
37. In the metaphor: LavaLamp is **\*what makes this castle yours\***. Without LavaLamp, the castle is generic; with LavaLamp, the castle has identity that cannot be transferred to another castle without the substrate also transferring.
38. Relevant LL entries: LL-001 (substrate-bound identity), LL-002 (visual/security decoupling), LL-003 (Lorenz-96 SDE choice), LL-006 (Lyapunov-spectrum residue audit), LL-007 (chaos-guard), LL-008 (resolution-bounded security claim), LL-018 (per-class adversary quantification), LL-019 (side-channel hardening), LL-021 (worst-case-adversary bound – the only LL-NNN entry currently at `:proved`). The five joint-defense meta-claims (LL-030 through LL-034) formalise structurally cohesive triads of these entries; see the LavaLamp whitepaper §5 for the joint-defense template.

### 2.2 PharOS – membrane checkpoint at the OS layer

-----

39. PharOS is the **\*\*OS-layer authentication membrane\*\*** – the checkpoint that filters identity claims into and out of the device's authentication subsystems. PharOS exposes LavaLamp's verifier API (LL-023) to the OS authentication infrastructure: PAM on Linux, Authorization Plug-in on macOS, Credential Provider on Windows. Each platform shim translates the verifier's accept / reject Bool into platform-native authentication results without leaking distance information to the threshold (LL-017 no-oracle preservation through the membrane).
40. In the metaphor: PharOS is the membrane gate. It does not block everything (the castle has to be useful); it filters selectively (the membrane is permeable to legitimate identity claims and impermeable to spoofed ones).
41. Relevant LL entries: LL-011 (registration ceremony), LL-012 (cold-start window), LL-013 (cross-config transitions),

LL-014 (threshold calibration), LL-017 (verification no-oracle), LL-023 (consumer-API surface). PharOS is currently forthcoming; the published whitepaper for it has not yet been released. The LL-023 surface in LavaLamp's spec carries the structural commitment that PharOS will instantiate.

## 2.3 Lazarus – backend / inner sanctum

---

- 42. Lazarus is the **\*\*substantive backend\*\*** – the inner sanctum where the work actually happens. The biblical-resurrection naming reflects the substance-under-the-hood / minimal-UI duality: what the user sees is utilitarian; what does the work is alive and persistent.
- 43. In the metaphor: Lazarus is the keep at the centre of the castle. The membrane (PharOS) checks identity at the gate; the keep (Lazarus) is where the high-value work happens with the assurance that the membrane is doing its job.
- 44. Lazarus has its own repository at [github.com/IridiumSoftware/lazarus](https://github.com/IridiumSoftware/lazarus). Cross-Triad documentation will land as PharOS matures and the umbrella becomes operationally meaningful.

## 2.4 The deployment-stack triple

---

- 45. LL-022 (downward – what LavaLamp depends on from the OS) + LL-023 (upward – what consumers depend on LavaLamp for) + LL-024 (operational – real-sensor instantiation strategy) form the **\*\*deployment-stack triple\*\***: a closure-of-three at the spec level that articulates the *\*operational preconditions\** for the immune-system architecture to work.
- 46. In the metaphor: the deployment-stack triple specifies the *\*environmental assumptions\** under which the castle's immune system functions. If those assumptions fail, the immune system cannot do its job – the castle is still standing but the alarm does not ring.
- 47. The deployment-stack triple is one of the five engine-discovered structurally-cohesive triads in LavaLamp's spec. The others (LL-016 + LL-024 + LL-029 sensor-defense, LL-007 + LL-019 + LL-022 reseed-oracle defense, LL-024 + LL-028 + LL-029 round-3 Tier 3, LL-006 + LL-023 + LL-029 status-tier-diverse, and LL-023 + LL-024 + LL-029 operational-deployment-stack) are formalised as joint-defense meta-claims LL-030 through LL-034. See the LavaLamp whitepaper §5 and THEOREMS.md theorems 18-35 for the structural composition shape.

---

## 3. WHAT THE TRIAD DEPLOYMENTS DO AND DON'T ADDRESS

---

### 3.1 Within scope

---

- 48. The Triad Deployments collectively address (at the primitive tier):

| Threat class

| Coverage

|

|                              |                                                                                        |
|------------------------------|----------------------------------------------------------------------------------------|
| A1 remote software           | Strong (LL-006, LL-007, LL-018)                                                        |
| A2 unprivileged user-space   | Strong (LL-006, LL-018)                                                                |
| A3 kernel-level              | OOS by declaration (LL-015)                                                            |
| A4 software side-channel     | Partial (timing yes, sensor-manipulation partial via cross-validation; LL-016, LL-019) |
| A5 registration-time insider | Partial (TPM defends substitution; LL-020 ε-DP defends observation)                    |
| A6 time-localized observer   | Partial (single-window observation bounded; LL-012 cold-start protocol)                |

49. Most of the strong-coverage cases are cockroach-class threats. The catapult-class threats partially covered (A4 / A5) require deployment-tier mitigations layered on the spec's primitive-tier defenses.

### 3.2 Out of scope at the primitive tier

50. The threat classes that the Triad Deployments do **\*\*not\*\*** directly defend against at the primitive tier:

| Threat class                               | Why OOS                                                     | Defensive layer                                                                              |
|--------------------------------------------|-------------------------------------------------------------|----------------------------------------------------------------------------------------------|
| A7 passive-emanation (catapult)            | Substrate radiates regardless of software discipline        | Hardware shielding (Faraday / TEMPEST); distance; EMI-filtered power. LL-025 declares scope. |
| Active substrate manipulation (catapult)   | Modified hardware attests to itself                         | Supply-chain integrity; trusted-vendor procurement; tamper-evident packaging.                |
| Software supply chain compromise (hybrid)  | Compromised dep chain produces compromised binary           | Reproducible builds; signed packages; SBOM; sigstore-like infra.                             |
| Social / governance / legal compulsion     | Substrate-bound ID does not help when the user is compelled | Operational: legal advice; jurisdictional considerations; cordon sanitaire.                  |
| Cryptographic primitive failure (catapult) | LavaLamp depends on TPM / RNG / PQC correctness             | Algorithm migration; PQ readiness; crypto defense-in-depth.                                  |
| Operational / human error (cockroach)      | Triad gives device identity, not user identity              | UEBA; training; MFA; phishing-resistant auth.                                                |
| Composability seams                        | Interfaces between systems are where                        | Integration testing; formal                                                                  |

|                                                |                                     |                                                                       |
|------------------------------------------------|-------------------------------------|-----------------------------------------------------------------------|
| (cockroach)                                    | attacks live                        | seam-conformance verification.                                        |
| Long-term substrate drift (cockroach patience) | Components age; replacements happen | Re-registration protocols; audit trails; deployment-context lifetime. |
| Quantum threat horizon (catapult)              | Q-Day breaks classical crypto       | PQC migration; QKD complementarity.                                   |

51. **\*\*Note the asymmetry.\*\*** Most of the OOS classes are catapult-class. Cockroach classes are mostly within scope or partially addressed. The catapult classes typically require *\*layered defenses\** at deployment time, not primitive-tier reinforcement.

### 3.3 Load-bearing assumptions

---

52. For the Triad Deployments to be operationally meaningful, the following assumptions must hold:

1. **\*\*Substrate not under physical interdiction.\*\*** No supply-chain compromise of the hardware before delivery.
2. **\*\*Substrate not under passive electromagnetic observation above noise floor.\*\*** TEMPEST-class observation either not occurring or bounded by deployment context (Faraday / distance / EMI-filtered).
3. **\*\*Software dependency graph integrity.\*\*** Build-tool and library dependencies are not tampered.
4. **\*\*Cryptographic primitives function as specified.\*\*** TPM, PQC, RNG behave per their standards.
5. **\*\*User is not compelled.\*\*** Substrate-bound identity does not survive legal / political coercion.
6. **\*\*Composability seams are individually intact.\*\*** LavaLamp ↔ PharOS ↔ OS-auth ↔ MFA integration testing is performed and conformance verified.
7. **\*\*Substrate is stable over the deployment's lifetime.\*\*** Or, if not, re-registration protocols are followed.
8. **\*\*Quantum threat horizon has not arrived for adjacent crypto.\*\*** Or, if it has, PQC migration is complete for the specific algorithms the protocol layer touches.

53. These are *\*load-bearing assumptions\**. Deployments at high threat tiers should treat each as an explicit operational requirement and verify it independently, not assume it.

---

## 4. DEFENSE-IN-DEPTH: COMPLEMENTARY SYSTEMS

---

54. The Triad Deployments are *\*necessary but not sufficient\** for high-assurance identity. Layered defenses are required at multiple tiers for the immune-system metaphor to work in practice.

### 4.1 Hardware-tier complements

---

55. - **\*\*Faraday-rated enclosure / TEMPEST-rated equipment\*\*** for

- passive-emanation defense.
- **\*\*EMI-filtered power conditioning\*\*** (consumer-grade for cockroach-tier; SCIF-grade for state-actor tier).
- **\*\*Shielded cables\*\*** (USB / HDMI / ethernet) for radiated-emission containment.
- **\*\*Supply-chain integrity:\*\*** trusted-vendor procurement; hardware bill-of-materials provenance; tamper-evident packaging.
- **\*\*TPM 2.0 / Secure Enclave / TrustZone\*\*** for hardware root-of-trust (per LL-022's required-mechanism enumeration).
- **\*\*Decommissioning protocols\*\*** for end-of-life (data-bearing components do not escape the security boundary).

#### 4.2 Software-tier complements

---

56. - **\*\*Reproducible builds\*\*** so binaries match published source.
- **\*\*Signed packages plus sigstore-like infrastructure\*\*** for supply-chain provenance.
  - **\*\*SBOM (Software Bill of Materials)\*\*** for dependency transparency.
  - **\*\*Lockfile discipline\*\*** (already in place for Julia, Lean, and other ecosystems in the canonical repositories).
  - **\*\*Static analysis\*\*** for known-vulnerability-pattern detection.
  - **\*\*Privilege separation\*\*** – multiple OS users / containers / VMs depending on threat tier.

#### 4.3 Operational-tier complements

---

57. - **\*\*MFA / phishing-resistant authentication\*\*** layered on top of Lavalamp's device identity (FIDO2; WebAuthn).
- **\*\*UEBA (User and Entity Behavior Analytics)\*\*** for cockroach-class detection at the user / device-pattern level.
  - **\*\*Security training\*\*** for human-facing attack surfaces.
  - **\*\*Incident response.\*\*** Assume compromise eventually happens; prepare to detect and recover rather than only prevent.
  - **\*\*Logging and forensic-trail discipline\*\*** for post-incident analysis.

#### 4.4 Strategic-tier complements

---

58. - **\*\*Cordon sanitaire under ruin-class risk:\*\*** structural disconnection with regulated membrane; resist taxonomy-greed; Taleb-asymmetric upside-vs-downside accounting.
- **\*\*Jurisdictional placement\*\*** considerations (legal-compulsion threat tier).
  - **\*\*Cryptographic-primitive monitoring\*\*** (track NIST PQC standardisation, TPM advisory feeds, etc.); migration planning baked into deployment lifecycle.
  - **\*\*Regime-awareness discipline.\*\*** Match threat-tier to defensive investment; don't over-defend low-tier deployments or under-defend high-tier ones.
59. **\*\*The defense-in-depth principle applies fractally.\*\*** Each tier has its own immune system at the appropriate scale and resolution. Hardware-tier complements provide the perimeter; software-tier complements provide the integrity guarantees; operational-tier complements provide the runtime monitoring; strategic-tier complements provide the macro-frame within

which the rest operates. None of these tiers is sufficient alone; the Triad Deployments occupy the substrate-tier immune-system role within this fractal stack.

---

## 5. SYMBIOSIS, NOT VICTORY

---

60. The honest summary of what the Triad Deployments are for:

- They are not invincible. They cannot defeat all attackers. That is the wrong claim shape.
- They raise the attacker's cost across multiple axes, formalised as five joint-defense meta-claims with machine-checked structural composition (LavaLamp THEOREMS.md §LL-030-LL-034).
- They make sub-state-level attacks economically unjustified relative to easier targets in the ecosystem.
- They make state-level attacks expensive AND detectable rather than impossible. Detection-postured, not prevention-postured.
- They depend on layered defenses at hardware, software, operational, and strategic tiers to provide high-assurance identity in adversarial deployments.
- They are sufficient at the primitive tier for most consumer use without additional layering, given operational-discipline hygiene.

61. **\*\*Symbiosis is the equilibrium goal.\*\*** The defender's job is to position the target so the attacker's cost-benefit math says "go elsewhere," then to detect and respond to the rare actors who attack anyway. The autopoietic-closure principle expressed at the predator-prey level: the closure is not "kill all enemies"; it is "maintain metabolic conditions under which most actors do not bother and we detect the rare ones who do."

62. The metaphor returns: we are alive and have identity. The castle has identity that cannot be transferred to another castle without the substrate also transferring. The membrane is selectively open. The immune system is constantly watching, waiting, digesting, incorporating, denying, obfuscating, repairing – at multiple scales and resolutions. It is impossible for the enemy to provide a persistent attack to beat us without dominating in scale or injecting a virus / cancer under our nose.

63. Both of those defeat conditions are state-actor-tier and rare. For everyone else – the population of would-be attackers from cockroach-tier opportunists to catapult-tier criminal organisations – the Triad Deployments shift the cost-benefit math toward "find easier food elsewhere," which is the same equilibrium the biological immune system maintains with its surrounding ecology.

64. That equilibrium is the goal. Not victory, not invincibility, not the elimination of all adversaries. **\*\*Symbiosis.\*\*** A federation of intelligences that mostly leave each other alone, with structural defenses at the substrate level for the rare cases where adversarial response is metabolically justified.

---

## 6. REFERENCES AND CANONICAL ARTEFACTS

---

### 65. Architectural and structural references:

Green, A. (2026). 'Possibilistic Security: Identity as Closure Residue on Ternary Causal Hypergraphs.' v1.1, May 2026. [github.com/IridiumSoftware/possibilistic-security](https://github.com/IridiumSoftware/possibilistic-security).

Green, A. (2026). 'Closure Forces Structure: The Standard Model from Rosen Closure on Ternary Causal Hypergraphs.' v1.4, May 2026. [github.com/IridiumSoftware/closure-forces-structure](https://github.com/IridiumSoftware/closure-forces-structure).

### 66. Per-deployment whitepapers:

Green, A. (2026). 'LavaLamp: Substrate-Bound Identity via Chaotic-SDE Residue Audit.' v1.0, May 2026. [github.com/IridiumSoftware/lavalamp](https://github.com/IridiumSoftware/lavalamp).

(Lazarus and PharOS whitepapers forthcoming.)

### 67. Canonical repositories:

- [github.com/IridiumSoftware/lavalamp](https://github.com/IridiumSoftware/lavalamp) – substrate-bound identity primitive (LavaLamp).
- [github.com/IridiumSoftware/lazarus](https://github.com/IridiumSoftware/lazarus) – security companion backend (Lazarus).
- [github.com/IridiumSoftware/possibilistic-security](https://github.com/IridiumSoftware/possibilistic-security) – architectural parent paper.
- [github.com/IridiumSoftware/closure-forces-structure](https://github.com/IridiumSoftware/closure-forces-structure) – physics foundation paper.
- PharOS: forthcoming.

### 68. The Triad Deployments umbrella does not currently have its own repository; it is a portfolio-level framing. An umbrella repository may be created when PharOS is closer to release and the cross-deployment infrastructure becomes operationally meaningful.

---